

## GROUP B

### EXPERIMENT: 1

#### WordCount.java

```
import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class WordCount {

    public static class TokenizerMapper
        extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
            ) throws IOException, InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken());
                context.write(word, one);
            }
        }

    }

    public static class IntSumReducer
        extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
            Context context
            ) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
```

```

job.setCombinerClass(IntSumReducer.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

**output: output file (part-r-00000)**



```

11 For 1
12 Gibson 1
13 Magazine. 1
14 Map 1
15 Project 1
16 Reduce 1
17 Shadow 2
18 Shadow, 1
19 Smith 1
20 Smith's 1
21 Story 1
22 Street 2
23 The 4
24 Walter 1
25 WordCount 1
26 a 4
27 about 1
28 an 1
29 and 4
30 be 1
31 broadcasts 1
32 by 2
33 can 1
34 contained 1
35 contains 1
36 count 1
37 counts 1
38 crime-fighting 1
39 data 1
40 dated 1
41 different 1
42 each 1
43 evil 1
44 example 1
45 figure 1
46 files 2
47 files, 1

```

Plain Text ▾ Tab Width: 8 ▾ Ln 22, Col 10 ▾ INS

## GROUP B

### EXPERIMENT: 2

#### Code:

#### 1> LogFileMapper.java (Use for mapping the IP addresses from input csv file)

```
package LogFileCountry;
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class LogFileMapper extends MapReduceBase implements Mapper<LongWritable, Text, Text,
IntWritable> {
    private final static IntWritable one = new IntWritable(1);

    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter
reporter) throws IOException {

        String valueString = value.toString();
        String[] SingleIpData = valueString.split("-");
        output.collect(new Text(SingleIpData[0]), one);
    }
}
```

#### 2> LogFileReducer.java (Use for reducing data received from mapper process to final output)

```
package LogFileCountry;
import java.io.IOException;
import java.util.*;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.*;

public class LogFileReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text,
IntWritable> {public void reduce(Text t_key, Iterator<IntWritable> values, OutputCollector<Text,IntWritable>
output, Reporter reporter) throws IOException {
    Text key = t_key;
    int frequencyForIp = 0;
    while (values.hasNext()) {
        // replace type of value with the actual type of our value
        IntWritable value = (IntWritable) values.next();
        frequencyForIp += value.get();
    }
    output.collect(key, new IntWritable(frequencyForIp));
}
}
```

### 3> LogFileCountryDriver.java (The driver code to run map-reduce on hdfs)

```
package LogFileCountry;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

public class LogFileCountryDriver {
    public static void main(String[] args) {
        JobClient my_client = new JobClient();
        // Create a configuration object for the job
        JobConf job_conf = new JobConf(LogFileCountryDriver.class);

        // Set a name of the Job
        job_conf.setJobName("LogFileIP");
        // Specify data type of output key and value
        job_conf.setOutputKeyClass(Text.class);
        job_conf.setOutputValueClass(IntWritable.class);
        // Specify names of Mapper and Reducer Class
        job_conf.setMapperClass(LogFileCountry.LogFileMapper.class);
        job_conf.setReducerClass(LogFileCountry.LogFileReducer.class);

        // Specify formats of the data type of Input and output
        job_conf.setInputFormat(TextInputFormat.class);
        job_conf.setOutputFormat(TextOutputFormat.class);

        // Set input and output directories using command line arguments,
        //arg[0] = name of input directory on HDFS, and arg[1] = name of output directory to be
        created to store the output file.

        FileInputFormat.setInputPaths(job_conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(job_conf, new Path(args[1]));

        my_client.setConf(job_conf);
        try { // Run the job
            JobClient.runJob(job_conf);
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}
```

### 4> log\_file.txt (Input file sample)

```
0.223.157.186 - - [15/Jul/2009:20:50:32 -0700] "GET /assets/js/the-associates.js HTTP/1.1" 304 -
10.223.157.186 - - [15/Jul/2009:20:50:33 -0700] "GET /assets/img/home-logo.png HTTP/1.1" 304 -
10.223.157.186 - - [15/Jul/2009:20:50:33 -0700] "GET /assets/img/dummy/primary-news-2.jpg HTTP/1.1" 304 -
10.223.157.186 - - [15/Jul/2009:20:50:33 -0700] "GET /assets/img/dummy/primary-news-1.jpg HTTP/1.1" 304 -
10.223.157.186 - - [15/Jul/2009:20:50:33 -0700] "GET /assets/img/home-media-block-placeholder.jpg HTTP/1.1" 304 -
-
10.223.157.186 - - [15/Jul/2009:20:50:33 -0700] "GET /assets/img/dummy/secondary-news-4.jpg HTTP/1.1" 304 -
10.223.157.186 - - [15/Jul/2009:20:50:33 -0700] "GET /assets/img/loading.gif HTTP/1.1" 304 -
10.223.157.186 - - [15/Jul/2009:20:50:33 -0700] "GET /assets/img/search-button.gif HTTP/1.1" 304 -
```

## 5> Output (part-00000.txt On Hadoop) (sample)

10.1.1.236	7
10.1.181.142	14
10.1.232.31	5
10.10.55.142	14
10.102.101.66	1
10.103.184.104	1
10.103.190.81	53
10.103.63.29	1
10.104.73.51	1
10.105.160.183	1
10.108.91.151	15
10.109.21.76	1
10.11.131.40	1
10.111.71.20	8
10.112.227.184	6
10.114.74.30	1
10.115.118.78	1
10.117.224.230	1



## Step For Logs File Code:

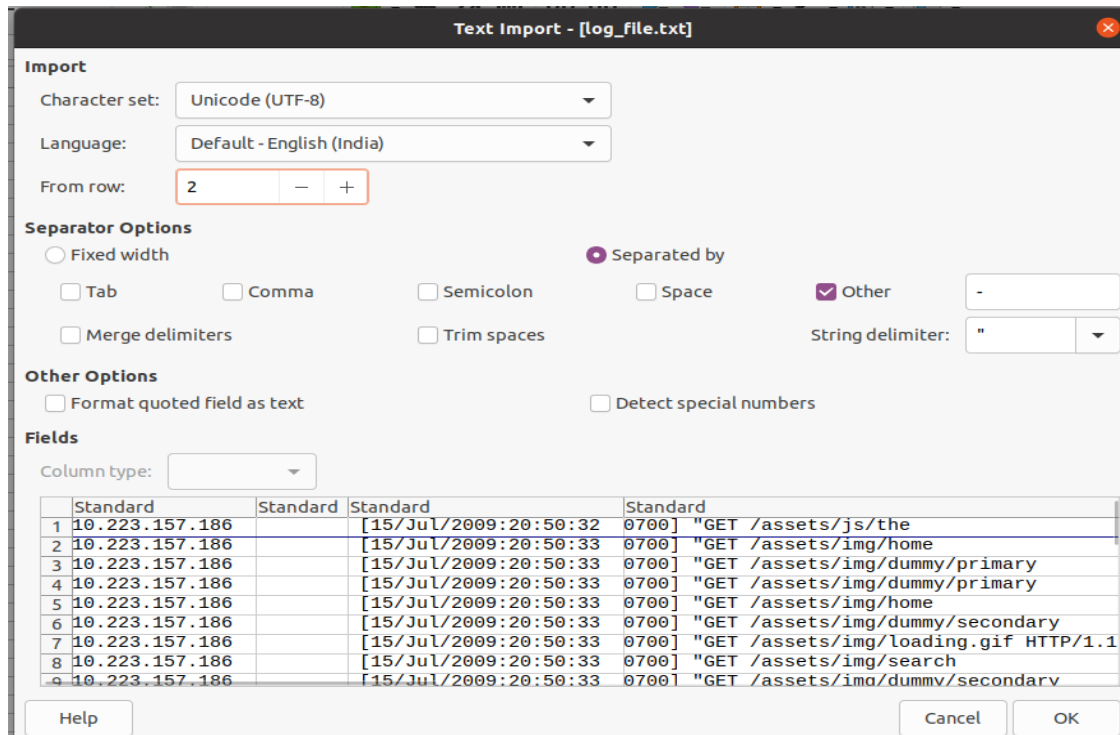
1. Starting Hadoop and check if it is started.

**\$ start-all.sh**

2. Create folder "LogFileTut". Copy the log\_file.txt given and create the java files.

- i. LogFileMapper.java
- ii. LogFileReducer.java
- iii. LogFileCountryDriver.java

3. Convert the log\_file.txt to .csv file. Open LibreOffice Calc-> Open -> log\_file.txt. Save As .csv in the LogFileTut folder.



4. Give Read permission to all the files in directories.

**\$ sudo chmod +r \*.\***

5. Set HADOOP\_CLASSPATH environment variable.

**\$ export HADOOP\_CLASSPATH=\$(hadoop classpath)**

or

**\$ export CLASSPATH=**

**"\$HADOOP\_HOME/share/hadoop/mapreduce/hadoop-mapreduce-client-core-3.2.2.jar:**

**\$HADOOP\_HOME/share/hadoop/mapreduce/hadoop-mapreduce-client-common-3.2.2.jar:**

**\$HADOOP\_HOME/share/hadoop/common/hadoop-common-3.2.2.jar: \$HADOOP\_HOME/lib/\*:**

**~/home/huser/Desktop/LogFileTut/\*"**

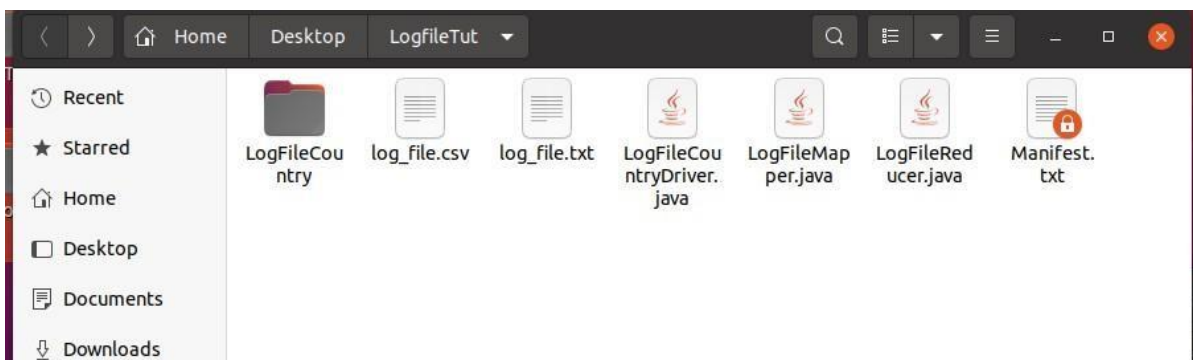

6. Compile the java code:

**\$ javac -classpath \$(HADOOP\_CLASSPATH) -d '/home/huser/Desktop/LogFileTut/exp\_classfile' \*.java**

7. Create Manifest.txt file.

```
huser@ubuntu-college:~/Desktop/LogfileTut$ javac -d '/home/huser/Desktop/LogfileTut/exp_classfile' LogFileMapper.java LogFileReducer.java LogFileCountryDriver.java
huser@ubuntu-college:~/Desktop/LogfileTut$ sudo gedit Manifest.txt
[sudo] password for huser:


(gedit:59507): Tepl-WARNING **: 22:24:16.402: GVfs metadata is not supported. Fallback to TeplMetadataManager. Either GVfs is not correctly installed or GVfs metadata are not supported on this platform. In the latter case, you should configure Tepl with --disable-gvfs-metadata.
```



8. Creation .jar file of classes:

**\$ jar -cvfm analyzelogs.jar Manifest.txt LogFileCountry/\*.class**

```
huser@ubuntu-college:~/Desktop/LogfileTut$ jar -cvfm analyzelogs.jar Manifest.txt LogFileCountry/*.class
added manifest
adding: LogFileCountry/LogFileCountryDriver.class(in = 1677) (out= 825)(deflated 50%)
adding: LogFileCountry/LogFileMapper.class(in = 1713) (out= 645)(deflated 62%)
adding: LogFileCountry/LogFileReducer.class(in = 1580) (out= 635)(deflated 59%)
```



9. Create a directory on HDFS .And check on localhost:9870

**\$ hdfs dfs -mkdir / LogFileExp**

**\$ hdfs dfs -mkdir / LogFileExp/Input**

**\$ hdfs dfs -mkdir / LogFileExp/Output**

10. Upload the log\_file.csv in hadoop dir /LogFileExp/Input



```
$ hdfs dfs -put '/home/huser/Desktop/LogFileTut/log_file.csv' /LogFileExp/Input
```

## 11. Running the jar file on Hadoop.

```
$ hadoop jar analyzlogs.jar /LogFileExp/Input /LogFileExp/Output
```

```
huser@ubuntu-college:~/Desktop/LogfileTut$ hadoop jar analyzlogs.jar /LogFileExp/Input /LogFileExp/Output
2022-04-12 22:51:25,988 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-04-12 22:51:27,403 INFO client.RMProxy: Connecting to ResourceManager at /127.0.0.1:8032
2022-04-12 22:51:35,208 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2022-04-12 22:51:36,276 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/huser/.staging/job_1649777619248_0001
2022-04-12 22:51:39,085 INFO mapred.FileInputFormat: Total input files to process : 1
2022-04-12 22:51:40,281 INFO mapreduce.JobSubmitter: number of splits:2
2022-04-12 22:51:40,821 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1649777619248_0001
2022-04-12 22:51:40,823 INFO mapreduce.JobSubmitter: Executing with tokens: []
2022-04-12 22:51:42,337 INFO conf.Configuration: resource-types.xml not found
2022-04-12 22:51:42,337 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2022-04-12 22:52:55,956 INFO impl.YarnClientImpl: Submitted application application_1649777619248_0001
2022-04-12 22:52:58,540 INFO mapreduce.Job: The url to track the job: http://ubuntu-college:8088/proxy/application_1649777619248_0001/
2022-04-12 22:52:58,584 INFO mapreduce.Job: Running job: job_1649777619248_0001
2022-04-12 22:57:02,946 INFO mapreduce.Job: Job job_1649777619248_0001 running in uber mode : false
2022-04-12 22:57:03,272 INFO mapreduce.Job: map 0% reduce 0%
2022-04-12 23:00:09,974 INFO mapreduce.Job: map 83% reduce 0%
2022-04-12 23:00:27,106 INFO mapreduce.Job: map 100% reduce 0%
2022-04-12 23:01:05,301 INFO mapreduce.Job: map 100% reduce 100%
2022-04-12 23:01:08,520 INFO mapreduce.Job: Job job_1649777619248_0001 completed successfully
2022-04-12 23:01:24,647 INFO mapreduce.Job: Counters: 54
```

## 12. Check the Output file.

```
$ hdfs dfs -cat /LogFileExp/Output/part-00000
```

```
huser@ubuntu-college:~/Desktop/LogfileTut$ hdfs dfs -cat /LogFileExp/Output/part-00000
10.1.1.236 7
10.1.181.142 14
10.1.232.31 5
10.10.55.142 14
10.102.101.66 1
10.103.184.104 1
10.103.190.81 53
10.103.63.29 1
10.104.73.51 1
10.105.160.183 1
10.108.91.151 1
10.109.21.76 1
10.11.131.40 1
10.111.71.20 8
```

The screenshot shows the Hadoop web interface with a modal window titled "File information - part-00000". The modal displays the following details:

- Block information: Block 0
- Block ID: 1073741897
- Block Pool ID: BP-1388353168-127.0.1.1-1647528100285
- Generation Stamp: 1073
- Size: 3838
- Availability: ubuntu-college

The "File contents" section shows a list of IP addresses and their corresponding counts:

IP Address	Count
10.240.170.50	1
10.241.107.75	1
10.241.9.187	1
10.243.51.109	5
10.244.166.195	5
10.245.208.15	20
10.246.151.162	3
10.247.111.104	9

## 13. Stop all processes :

```
$ stop-all.sh
```

**Conclusion:** Thus, we successfully implement, distributed application using MapReduce which processes a log file of a system.



## GROUP B

### EXPERIMENT: 3

#### Weather.java

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;
import java.util.StringTokenizer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.KeyValueTextInputFormat;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.util.*;
/**
 * This is an Hadoop Map/Reduce application for Working on weather data It reads
 * the text input files, breaks each line into stations weather data and finds
 * average for temperature, dew point , wind speed. The output is a locally
 * sorted list of stations and its 12 attribute vector of average temp , dew ,
 * wind speed of 4 sections for each month.
 */
public class Weather extends Configured implements Tool {
    final long DEFAULT_SPLIT_SIZE = 128 * 1024 * 1024;
    /**
     * Map Class for Job 1
     * For each line of input, emits key value pair with
     * station_yearmonth_sectionno as key and 3 attribute vector with
     * temperature , dew point , wind speed as value. Map method will strip the
     * day and hour from field and replace it with section no (
     * <b>station_yearmonth_sectionno</b>, <b>temperature,dew point , wind
```

$\ast/$

```

        }else{
            value_str = value_str.concat(str).concat(" ");
            break;
        }
    case 12://Wind speed
        if(str.equals("999.9")){//Ignoring rows wind speed is all 9
            skip = true;
            break loop;
        }else{ value_str = value_str.concat(str).concat(" ");
            break;
        }
    default: break;
    }
    counter++;
}
if(!skip){
    word.set(key_out);
    values.set(value_str);
    output.collect(word, values);
}
}
}
}
/**
 * Reducer Class for Job 1
 * A reducer class that just emits 3 attribute vector with average
 * temperature , dew point , wind speed for each of the section of the month for each input
 */
public static class Reduce extends MapReduceBase
    implements Reducer<Text, Text, Text, Text> {
    private Text value_out_text = new Text();
    public void reduce(Text key, Iterator<Text> values,
        OutputCollector<Text, Text> output, Reporter reporter) throws IOException {
        double sum_temp = 0;
        double sum_dew = 0;
        double sum_wind = 0;
        int count = 0;

        while (values.hasNext()) {
            String str = values.next().toString();
            StringTokenizer itr = new StringTokenizer(str);
            int count_vector = 0;
            while (itr.hasMoreTokens()) {
                String nextToken = itr.nextToken(" ");
                if(count_vector==0){
                    sum_temp += Double.valueOf(nextToken);
                }
                if(count_vector==1){
                    sum_dew += Double.valueOf(nextToken);
                }
                if(count_vector==2){
                    sum_wind += Double.valueOf(nextToken);
                }
            }
        }
    }
}

```

```

        count_vector++;
    }
    count++;
}
double avg_tmp = sum_temp / count;
double avg_dew = sum_dew / count;
double avg_wind = sum_wind / count;
System.out.println(key.toString()+" count is "+count+" sum of temp is "+sum_temp+" sum of
dew is "+sum_dew+" sum of wind is "+sum_wind+"\n");
String value_out = String.valueOf(avg_tmp).concat("
").concat(String.valueOf(avg_dew)).concat(" ").concat(String.valueOf(avg_wind));
value_out_text.set(value_out);
output.collect(key, value_out_text);
}
}
static int printUsage() {
    System.out.println("weather [-m <maps>] [-r <reduces>] <job_1 input> <job_1 output> <job_2
output>");
    ToolRunner.printGenericCommandUsage(System.out);
    return -1;
}
/**
 * The main driver for weather map/reduce program.
 * Invoke this method to submit the map/reduce job.
 * @throws IOException When there is communication problems with the job tracker.
 */
public int run(String[] args) throws Exception {
    Configuration config = getConf();
    // We need to lower input block size by factor of two
    JobConf conf = new JobConf(config, Weather.class);
    conf.setJobName("Weather Job1");

    // the keys are words (strings)
    conf.setOutputKeyClass(Text.class);
    // the values are counts (ints)
    conf.setOutputValueClass(Text.class);

    conf.setMapOutputKeyClass(Text.class);
    conf.setMapOutputValueClass(Text.class);
    conf.setMapperClass(MapClass.class);
    //conf.setCombinerClass(Combiner.class);
    conf.setReducerClass(Reduce.class);
    List<String> other_args = new ArrayList<String>();
    for(int i=0; i < args.length; ++i) {
        try {
            if ("-m".equals(args[i])) {
                conf.setNumMapTasks(Integer.parseInt(args[++i]));
            } else if ("-r".equals(args[i])) {
                conf.setNumReduceTasks(Integer.parseInt(args[++i]));
            } else {
                other_args.add(args[i]);
            }
        }
    }
}

```

```

        } catch (NumberFormatException except) {
            System.out.println("ERROR: Integer expected instead of " + args[i]);
            return printUsage();
        } catch (ArrayIndexOutOfBoundsException except) {
            System.out.println("ERROR: Required parameter missing from " +
                               args[i-1]);
            return printUsage();
        }
    }
    // Make sure there are exactly 2 parameters left.
    FileInputFormat.setInputPaths(conf, other_args.get(0));
    FileOutputFormat.setOutputPath(conf, new Path(other_args.get(1)));
    JobClient.runJob(conf);
    return 0;
}
public static void main(String[] args) throws Exception {
    int res = ToolRunner.run(new Configuration(), new Weather(), args);
    System.exit(res);
}
}

```

### Input: sample\_weather.txt (sample)

```

690190 13910 20060201_0 51.75 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_1 54.74 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_2 50.59 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_3 51.67 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_4 65.67 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_5 55.37 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_6 49.26 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_7 55.44 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000
690190 13910 20060201_8 64.05 33.0 24 1006.3 24 943.9 24 15.0 24 10.7 24 22.0 28.9 0.001 999.9 000000

```

### Output: part-00000.txt (on Hadoop)

```

690190_02_section1 53.87166666666666 25.899999999999995 7.774999999999998
690190_02_section2 54.761250000000001 25.900000000000006 7.774999999999999
690190_02_section3 53.250416666666667 25.899999999999995 7.774999999999996
690190_02_section4 52.447083333333333 25.900000000000006 7.774999999999999

```

### Weather Data Analysis Steps to run:

14. Starting Hadoop

**\$ start-all.sh**

15. Made A folder “WeatherAssi” and write Weather.java code.

28. Create new folder for input data.

29. Add input text file in the input data folder.

30. Create new folder to hold java class files.

31. Set HADOOP\_CLASSPATH environment variable.

**\$ export HADOOP\_CLASSPATH=\$(hadoop classpath)**

32. Create a directory on HDFS

```
$ hdfs dfs -mkdir /WeatherTut
```

```
$ hdfs dfs -mkdir /WeatherTut/Input
```

33. Checking on localhost:9870

34. Upload the input file (device) to that directory.

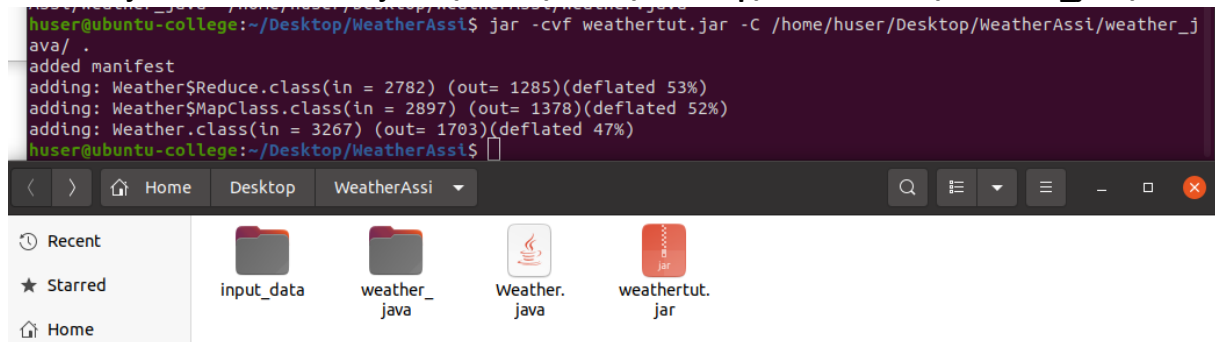
```
$ hdfs dfs -put input_data/sample_wheater.txt /WeatherTut/Input
```

35. Compile the java code:

```
$ javac -classpath $(HADOOP_CLASSPATH) -d '/home/huser/Desktop/WeatherAssi/
weather_java' /home/huser/Desktop/WeatherAssi/Weather.java
```

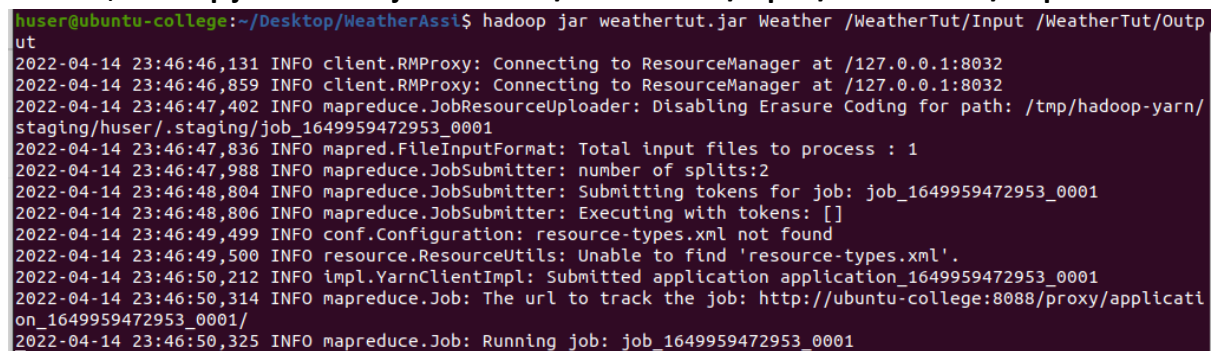
36. Creation .jar file of classes:

```
$ jar -cvf weathertut.jar -C /home/huser/Desktop/WeatherAssi/weather_java/ .
```

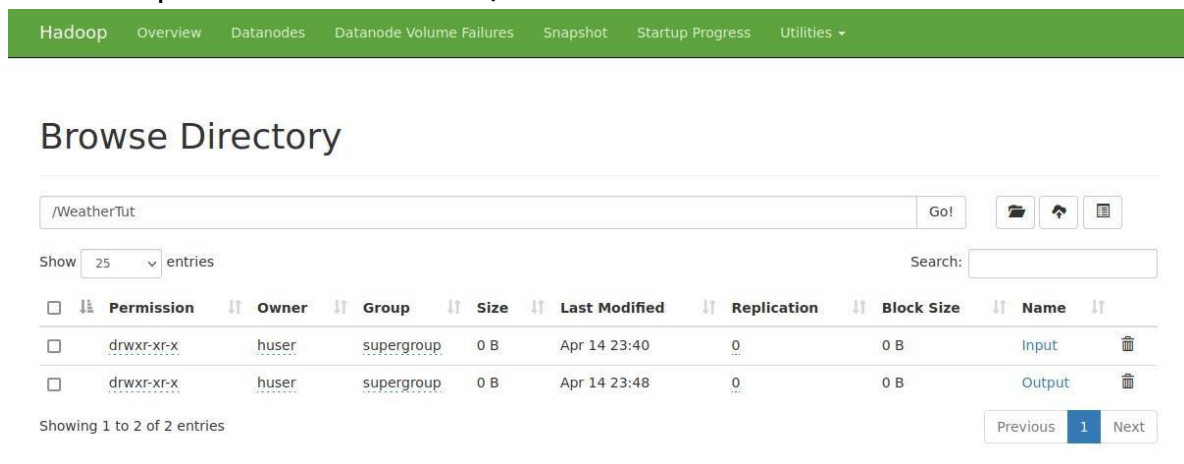


37. Running the jar file on Hadoop

```
$ hadoop jar weather.jar Weather /WeatherTut/Input /WeatherTut/Output
```



38. Check output on localhost:9870 /localhost:50070





## Browse Directory

/WeatherTut/Output Go! 📁 📄 📄

Show  entries Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
<input type="checkbox"/>	-rw-r--r--	huser	supergroup	0 B	Apr 14 23:48	1	128 MB	_SUCCESS
<input type="checkbox"/>	-rw-r--r--	huser	supergroup	296 B	Apr 14 23:48	1	128 MB	part-00000

Showing 1 to 2 of 2 entries Previous 1 Next

Overview Datanodes Datanode Volume Failures Snapshot Startup Progress Utilities

### Browse Directory

/WeatherTut/Output

entries

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	huser	supergroup	0 B	Apr 14 23:48	1	128 MB	_SUCCESS
-rw-r--r--	huser	supergroup	296 B	Apr 14 23:48	1	128 MB	part-00000

Showing 1 to 2 of 2 entries

File information - part-00000

[Download](#) [Head the file \(first 32K\)](#) [Tail the file \(last 32K\)](#)

Block information --

Block ID: 1073741918  
Block Pool ID: BP-1388353168-127.0.1.1-1647528100285  
Generation Stamp: 1094  
Size: 296  
Availability:  
• ubuntu-college

File contents

```
690190_02_section1 53.87166666666666 25.899999999999995 7.774999999999998
690190_02_section2 54.76125000000001 25.900000000000006 7.774999999999999
690190_02_section3 53.25041666666667 25.899999999999995 7.774999999999996
690190_02_section4 52.44708333333333 25.900000000000006 7.774999999999999
```

Close

39. Stop Hadoop services:

**\$ stop-all.sh**

**Conclusion:** Thus, we successfully, Locate dataset (e.g., sample\_weather.txt) for working on weather data which reads the text input files and finds average for temperature, dew point and wind speed.